

Final audit adjudication and exact correction plan

MathGov/RippleLogic v13.0 · MG-RL-13.0-20260923-RELEASE-G · 23 September 2026

Decision

The authorized targeted corrections are integrated into the governing sources. The cascade, seven scopes/dimensions, normalized RLS, rights non-compensation, CVaR, SGP split and component editions remain unchanged. This is a new exact correction build, not the prior frozen bytes under a reused identity. The original publication ZIP is preserved in the Complete archive.

Input ZIP SHA-256: 9f1865a69533e82725d90d907e920236ff462f76160b74ee455feddba9345e3f.
Reviewer file: Feedback from Opus 5.5 on v13.0 core 15 files.docx, SHA-256
00180594e83a60fdb95e53baa05ecdff0496a160a193bde8bb682a36ee8f48ad4. Findings were checked against the actual Office bytes rather than assumed from reviewer summaries. The report's limited-access qualifications were respected.

The two consequential items are recovery-mode continuity and compatible joint robustness testing. They were fixed in the Agent/Canon masters, not merely via a package note. Low-risk errata were corrected only where reproduced. The remaining proposals were retained, narrowed or deferred as described below.

G-01 · B-01 · IMPLEMENTED

Where: Agent §43.1 and §43.4

Before / concern: The old recovery text allowed resumption when the backend returned or resumption in the prior mode, despite §8 and §23.7 restrictions.

Change / decision: Backend/network loss enters MODE 0; recovery reconciles logs, configuration, current authority and affected qualification. Section 8 re-entry, including authenticated operator approval, must complete before elevated mode resumes. External protective interlocks retain only their own existing authority.

Why: Closes an operationally material bypass without weakening the existing permission ladder.

Evidence: Synthetic recovery tests cover each missing precondition, stale authority, repeated outage and complete re-entry. Existing invariant and successor tests remain.

G-02 · B-02 · IMPLEMENTED WITH EXPLICIT PREMISES

Where: Canon §10.4B; R.38A; reference fixture

Before / concern: The earlier text required uncertainty and dependence variants but did not state their compatible joint evaluation. R.38A lacked a specific cross-option disposition.

Change / decision: Apply required uncertainty-scale variants within every required dependence treatment. Evaluate other compatible, plausibly coexisting, material joint stresses without double-counting or crossing incompatible models. R.38A explicitly stipulates independent symmetric option-level generators with within-option perfect dependence and covariance zero by construction.

Why: Prevents separately passing tests from concealing a failing combination. The fixture is a stipulated mathematical positive control, not observed evidence or a convenient empirical non-trigger.

Evidence: Independently reproduced adverse $\rho=-1$ and $\text{scale}\times 2$ failures: A-B 1.87134858466; A-C 1.99750467776. Declared independent fixture retains 2.64135271898 and 2.76685785546. Removing its warranted premise blocks the scoped unique claim.

G-03 · A-01 · IMPLEMENTED

Where: Canon Appendix P.8

Before / concern: The numerator pointed to nonexistent Table P-8.

Change / decision: Reference Table P-4.

Why: Repairs navigation to the actual source values.

Evidence: Exact source text and table identity checked.

G-04 · A-02 · IMPLEMENTED

Where: Canon Appendix P.3

Before / concern: P.3 omitted confidence while P.8 reconstructed the same values using $c_k=0.85$.

Change / decision: State $c_k=0.85$, reach and adjustment 1; forward construction includes confidence; inverse divides by $\text{beta}\cdot\text{tau}\cdot\text{likelihood}\cdot\text{confidence}$.

Why: Makes the declared forward and inverse models agree without changing the supplied impact table or RLS result.

Evidence: Full-precision replay independently reconstructs every declared nonzero input under the clarified convention.

G-05 · A-03 · IMPLEMENTED

Where: Workbook RLS!A35; Config!C22; RLS!F67

Before / concern: Raw score difference was called Gap and an adjacent nominal label could be read as a final decisive result.

Change / decision: Use “ $\Delta\text{RLS (raw)}$ ”, “ $\text{SignedGap} > \delta$ in every required comparison”, and “Nominal demo only; see final verdict”.

Why: Separates raw difference, normalized discrimination and final robustness-qualified verdict. Concise labels fit existing cells.

Evidence: Formula/cache fingerprint unchanged; native recalculation and cold reopen compare all 20,212 formula results.

G-06 · A-06 · IMPLEMENTED

Where: Canon §13.3A and §13.5

Before / concern: Near-floor supermajority language did not locally clarify Tier 3 scope or final blended weight versus raw ballot share.

Change / decision: Specify Tier 3 HDW and the proposed final combined weight after the declared HDW construction. Preserve the 0.02 floor margin, 2/3 approval and stricter applicable rules.

Why: Resolves two plausible implementation readings; does not extend this rule to PLSS-only profiles.

Evidence: Exact clauses checked; existing floor, HDW and PLSS tests retained.

G-07 · B-03 · IMPLEMENTED

Where: Canon P.8 and R.19.7

Before / concern: Some printed derived values were based on already-rounded intermediate columns.

Change / decision: P.8 $Q=20886/30625 \approx 0.681991836735$; R.19.7 $\sigma=0.000502303633$; U6D5 x displays 0.017652. State the unrounded calculation basis and rounded display distinction.

Why: Uses one reproducible precision convention without changing the underlying model or final six-decimal RLS.

Evidence: Independent rational-weight replay: P $RLS(B)=0.0066483769031887375$; P $\sigma \approx 0.000951728456704$; R.19 $\sigma \approx 0.000502303632743$.

G-08 · B-04 · IMPLEMENTED

Where: Canon P.7 skipped-stage table and immediately adjacent explanation

Before / concern: The table used an inconsistent prior-elimination token and described A as eliminated by both NCRC and TRC, unlike P.9 short-circuit semantics.

Change / decision: Use CSV_NOT_EVALUATED_AFTER_PRIOR_FAILURE. State elimination by NCRC and label the separate TRC computation audit-only.

Why: Prevents contradictory canonical run states while retaining the valid counterfactual tail arithmetic.

Evidence: Token and surrounding explanation checked against P.9 and current state crosswalk.

G-09 · A-09 · IMPLEMENTED AS FIXTURE CLARIFICATION

Where: Canon P.4 and R.19.7

Before / concern: Method B examples displayed zero uncertainty on synthetic zero cells without a sufficiently local exact-input stipulation.

Change / decision: State these exact synthetic zeros are exact only for the proxy demonstration. Empirical assessed zero remains subject to measurement/coding/model resolution and transformation consistency.

Why: Clarifies the declared example instead of adding an arbitrary uncertainty floor or confusing missing evidence with zero.

Evidence: Method B equation unchanged. Workbook's separate all-49-cell uncertainty basis unchanged. Known-zero/NE and unknown propagation tests retained.

G-10 · B-05 · IMPLEMENTED

Where: Workbook Audit_Flags!C29:C31 and A39

Before / concern: Three Agent-owned runtime tokens were displayed with PCC-like severities and missing local ownership explanation.

Change / decision: Use RUNTIME_BLOCK, RUNTIME_BLOCK and RUNTIME_INVALID; explain runtime ownership and that RUNTIME_* does not extend the PCC vocabulary. Flags remain frozen NO records.

Why: Corrects a namespace/serialization ambiguity without building a new flag engine or changing a verdict.

Evidence: Actual values, unchanged formulas/caches and matching snapshot records checked; native results unchanged.

G-11 · Required synchronization · IMPLEMENTED

Where: All 14 terminal release appendices; workbook release cells/snapshots; manifests; reading copies; tests

Before / concern: Patched files require a distinct exact identity and current evidence, not reuse of the previous freeze claim.

Change / decision: Record new Build G with unchanged component editions. Update ten matching literal snapshots, static TOC page labels, document statistics, current guides, source indexes and manifests. Preserve all other inherited formatting.

Why: Prevents mixed-build release records and false staleness while keeping previous frozen artifacts immutable in provenance.

Evidence: All 336 table-format structures and section geometry preserved; 95-sheet workbook geometry preserved; all formula/cache records unchanged; 125 TOC references verified.

G-12 · A-04 / B-06 · RETAINED

Where: Workbook range guards and frozen rights rows

Before / concern: The review's native mutations show extra rows and subgroup edits produce STALE; individual rights records are intentionally frozen.

Change / decision: No SUMIFS range rewrite or live-rights evaluator added.

Why: Changing a correctly bounded frozen example into a different operational workbook would exceed the task.

Evidence: Existing mutation, staleness, input-validation and native acceptance suites retained.

G-13 · A-08 / A-13 / B-07 / B-08 · NO NEW CORE MECHANISM

Where: Local explanatory terms; external operating prompt; catastrophe and wrapper distinctions

Before / concern: Some observations concern external project instructions or already-governed context, not an independently reproduced Core bypass.

Change / decision: Do not mint new Canon tokens from local narrative labels, mutate external account memory, or change CVaR/waiver thresholds. Current README uses the actual cascade and

the existing claim boundary.

Why: Avoids architecture expansion and unwarranted general “catastrophic risks never averaged” statements.

Evidence: Scope and controls remain in the governing sources. No external prompt/settings or live publication changed.

G-14 · A-05 / A-07 / A-10 / A-11 / A-12 · DEFERRED OR COVERED BY GUIDANCE

Where: Nonblocking layout/navigation/legacy-label observations

Before / concern: These observations do not demonstrate a failed result in the current declared scope.

Change / decision: Do not move already-correct sections, add a new workbook coverage field or rebuild TOCs solely to expand a cleanup pass. Current reading guide identifies labels, zero basis, stream distinctions and manifest authority.

Why: Limits regression risk. Appendix RELEASE remains reachable through the reading index/bookmarks and document end even though it is not added to the numbered opening TOC.

Evidence: Current link/anchor checks and source parity run; no unresolved reported outcome-changing defect is hidden by this disposition.

Exact records and preservation

Verification/Exact_Document_Changes.json records 60 sequential targeted Word operations, including 42 terminal identity/evidence synchronizations across 14 documents.

Navigation_Changes.json records three changed Canon page labels.

Document_Statistics_Changes.json records necessary stored-statistics changes. No Word style, numbering, media or section-geometry parts changed.

Exact_Workbook_Changes.json records 20 literal-cell changes: ten principal edits and ten corresponding literal-snapshot mirrors. No row/column was inserted, no formula/cached result changed, and the original styles/geometry were preserved. A separately trusted current manifest identifies these new bytes. Reference_Code_Changes.json and the accompanying diff record test, fixture and identity-binding updates.

Evidence and claims

Read Verification_and_Readiness.md for actual executed evidence. Passing synthetic controls and arithmetic does not establish empirical calibration, complete external-registry conformance, lawful authority or deployment safety. R.38A’s new independent-generator premise is synthetic and declared; it is not evidence inferred from the absence of an observed covariance. Future real cases need their own warranted joint model or bound.